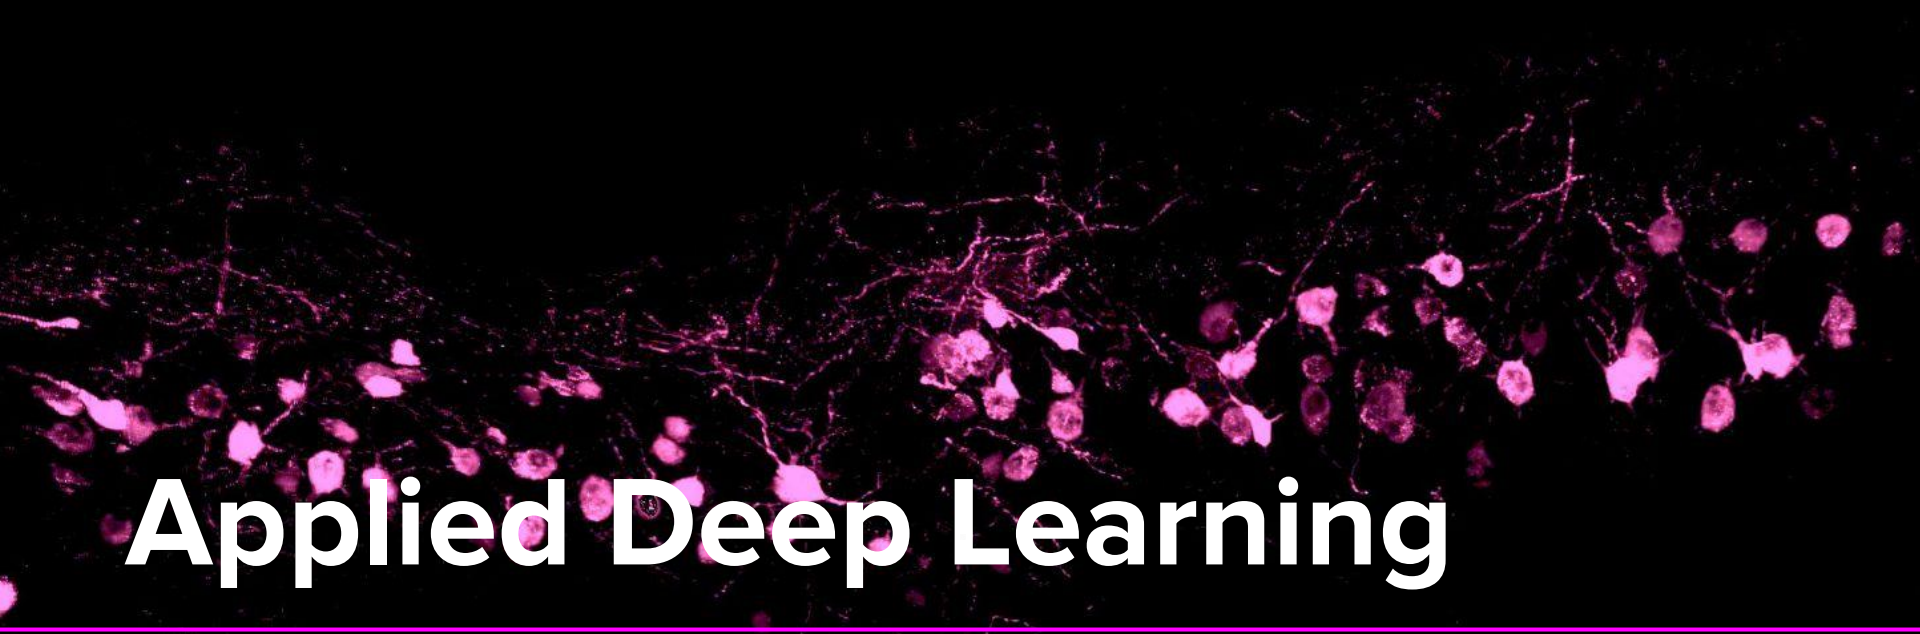

Applied Deep Learning

Introduction to Deep Learning
2025 / 2026

Alexander Pacha - TU Wien

Is this spam?

Congratulations! You have won €650.000,00 in the monthly Euro Millions / Google Promo lottery on April 26th 2019.

Please consult timothey.berrel@my-email.com for claiming your winnings with the following information:

1. Full Name
2. Address
3. Birth Date
4. Credit Card Number
5. CVC Checksum

Robert Avthantilayn
Online Coordinator

Is this spam?

Greetings,

I am Mrs. Margaret Ko May-Yee, Deputy Chairman and Managing Director of the Chong Hing Bank Limited. I write briefly to seek your collaboration in a multi-million transaction with good return for us on participation. Reply for details.

Regards,
Margaret Ko May-Yee

What do you see?

[Fig-2]

What do you see?

[Fig-3]

Machine Learning as Alternative

- Classical Programming quickly reaches its limits in certain situations
- Machine Learning can be a viable alternative to imperative programming

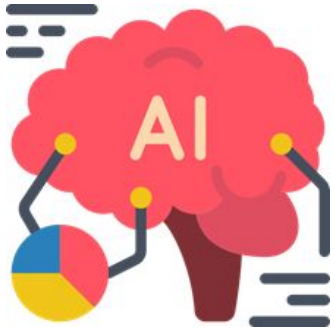

Machine Learning = Use data (Training) to answer questions (Inference)

Terminology

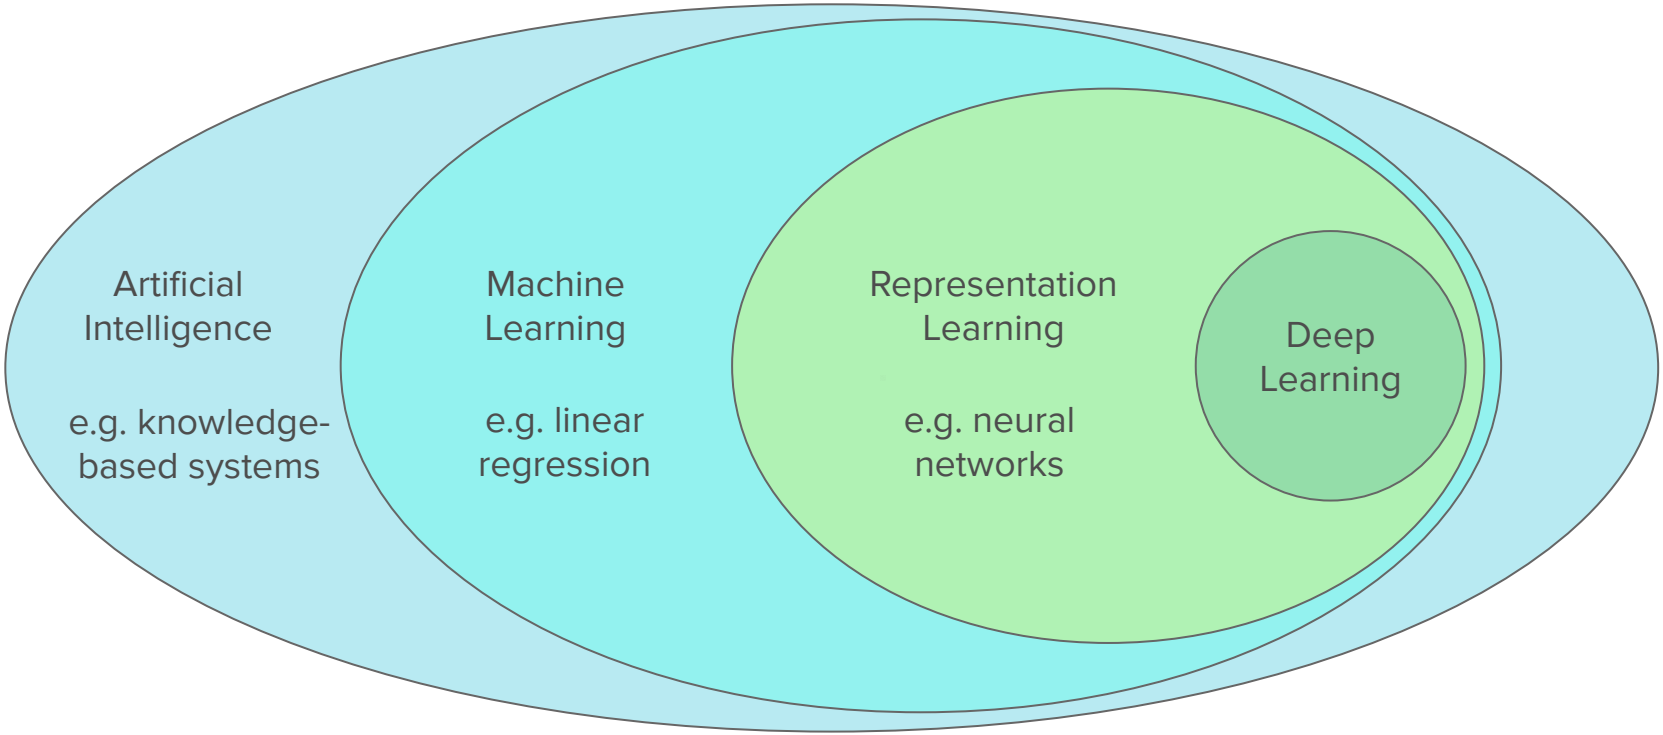

Meaningful Representation

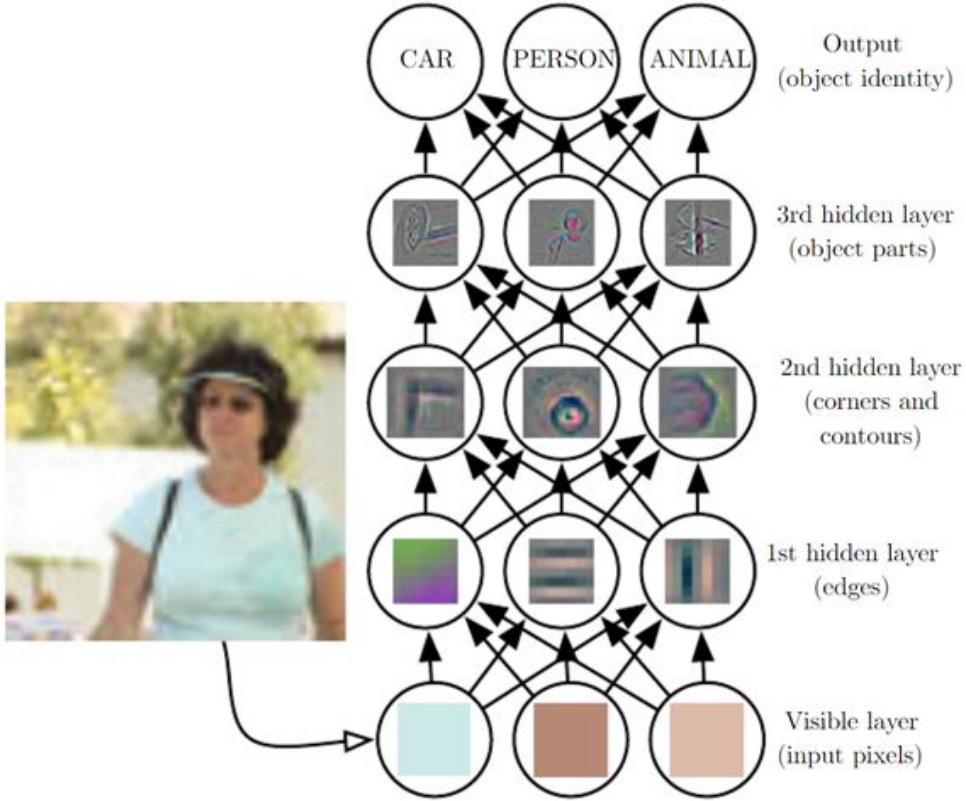

[1]

How Can a Machine Learn?

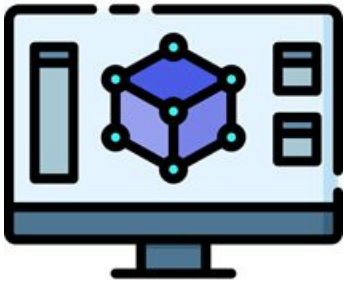

Model

Loss function

Optimization function

Image Classification

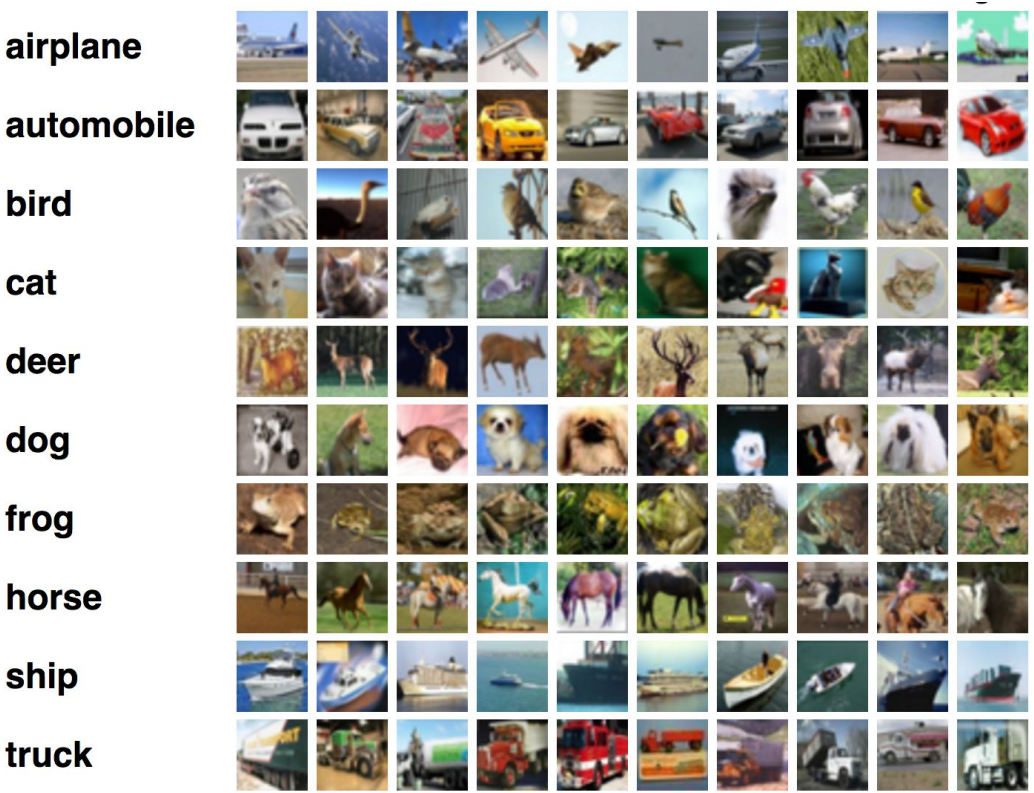

Image Classification

Melanoma

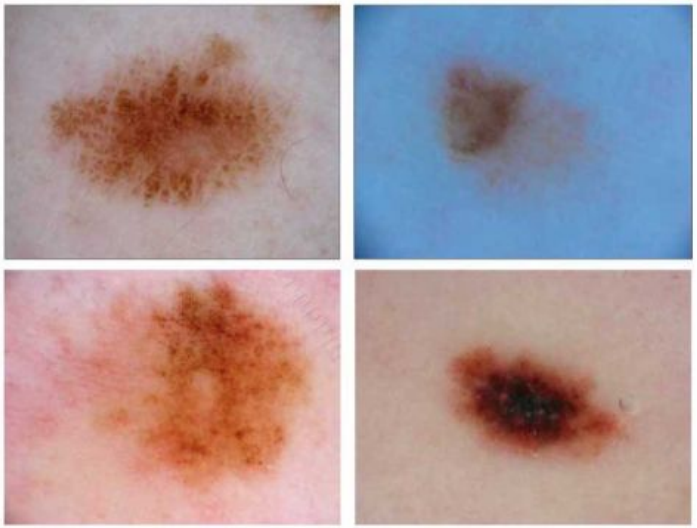

Benign

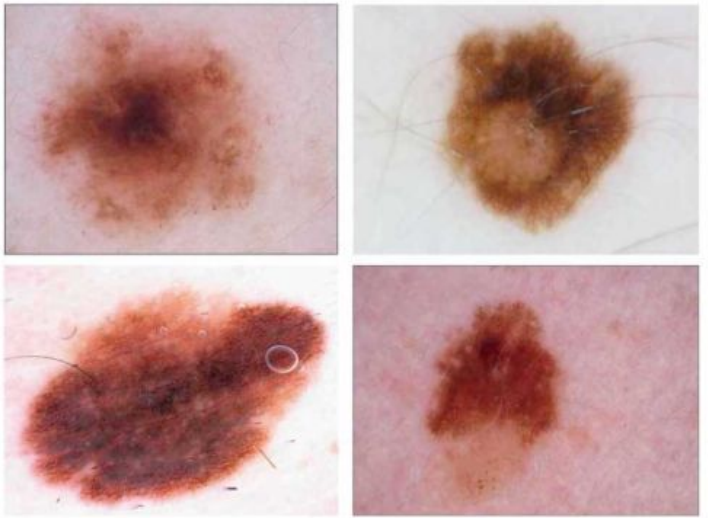

[5, 6]

Object Detection

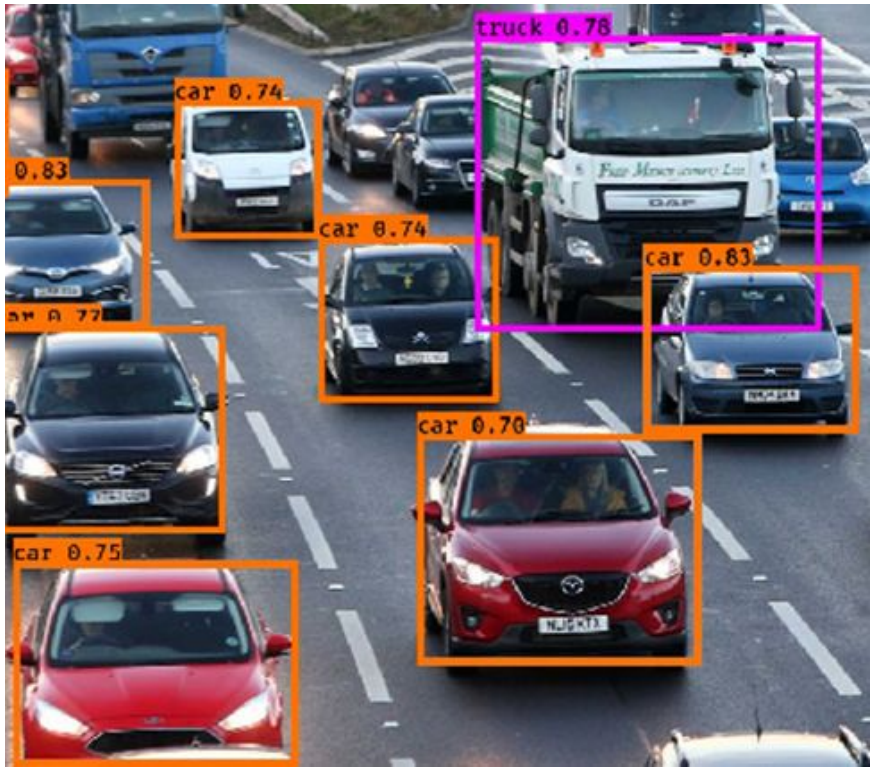

Voice Recognition and Natural Conversation

Google Duplex calls a hairdresser

Robo Calls

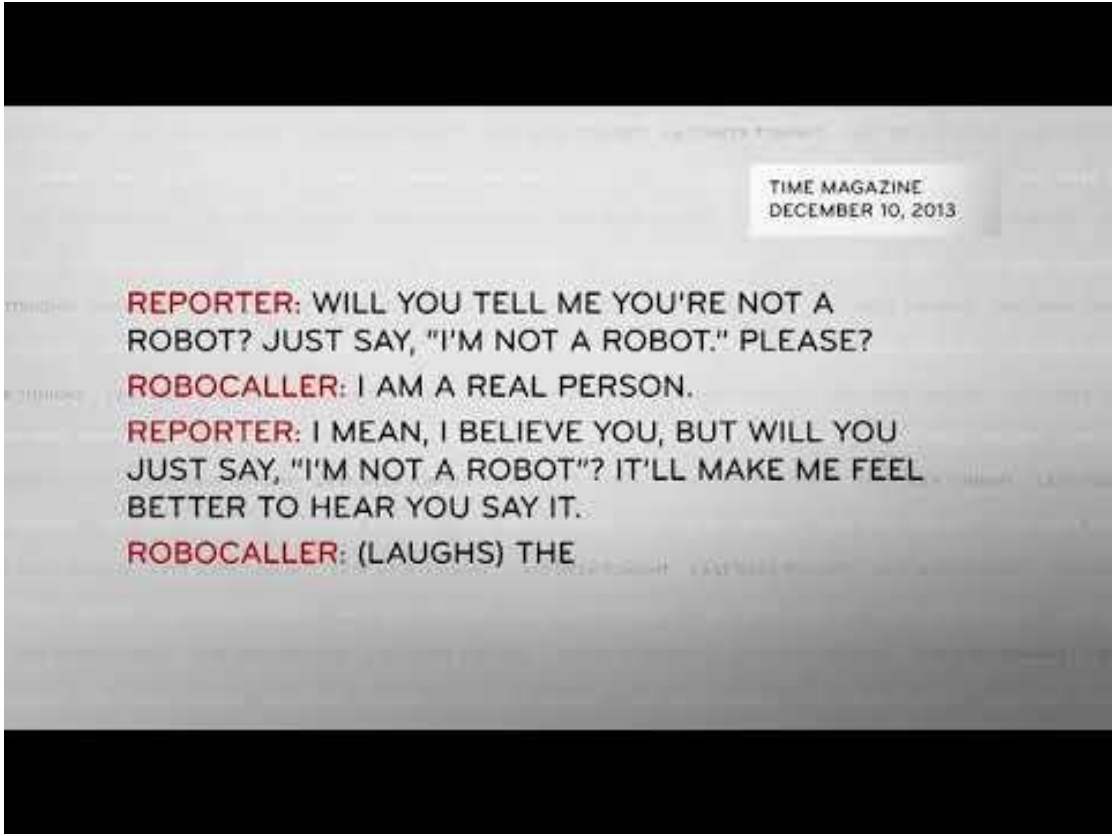
TIME MAGAZINE
DECEMBER 10, 2013
REPORTER: WILL YOU TELL ME YOU'RE NOT A ROBOT? JUST SAY, "I'M NOT A ROBOT." PLEASE?
ROBOCALLER: I AM A REAL PERSON.
REPORTER: I MEAN, I BELIEVE YOU, BUT WILL YOU JUST SAY, "I'M NOT A ROBOT"? IT'LL MAKE ME FEEL BETTER TO HEAR YOU SAY IT.
ROBOCALLER: (LAUGHS) THE

Music Harmonization

DeepBach

Creating new music

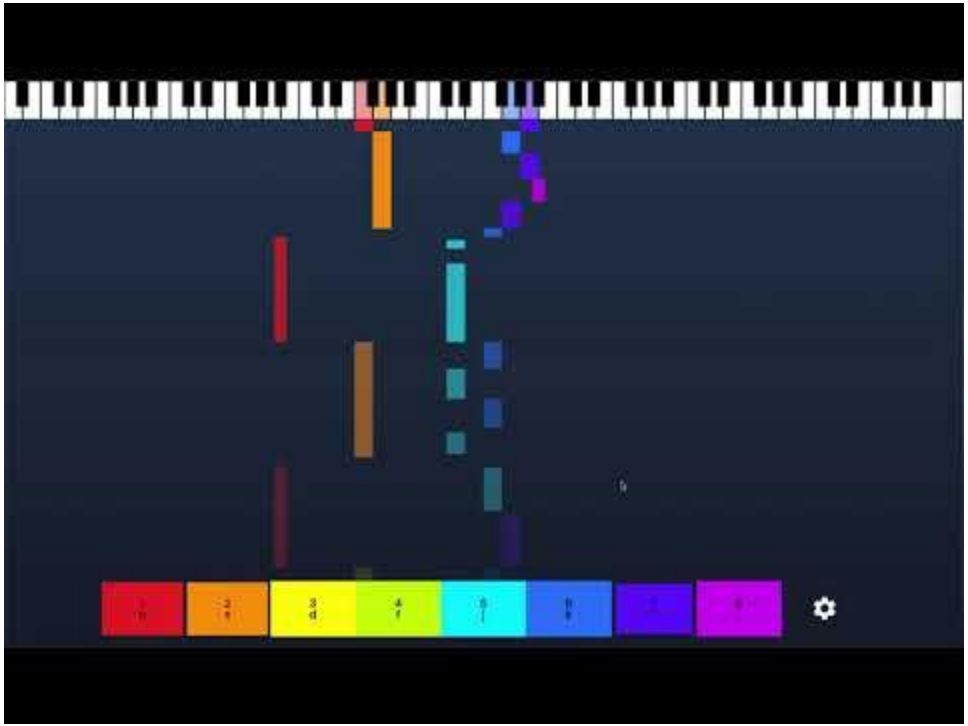

Creating new music

Creating Images from Text

Text Prompt: `an illustration of a baby fox in a suit playing guitar`

AI generated images:

[13, 14]

Creating Images from Text

Midjourney generations over time: “a hyper-realistic image of Harry Potter”
Source: Midjourney, 2024

Figure 2.3.6

Creating Videos from Text

Text Prompt: `A medium shot, historical adventure setting: Warm lamplight illuminates a cartographer in a cluttered study, poring over an ancient, sprawling map spread across a large table. Cartographer: "According to this old sea chart, the lost island isn't myth! We must prepare an expedition immediately!"`

OpenAI writes Python Code

```
    items: List[Item]

    def compute_total_price(self, palindrome_discount=0.2):
        """
        Compute the total price and return it.
        Apply a discount to items whose names are palindromes.
        """
        total_price = 0
        for item in self.items:
            if is_palindrome(item.name):
                total_price += item.price * (1 - palindrome_discount)
            else:
                total_price += item.price
        return total_price

    def print_receipt(self):
        """Print the total price and the price of each item"""
        print("Total price: {:.2f}".format(self.compute_total_price()))
        for item in self.items:
            print("{:>10}: {:.2f}".format(item.name, item.price))
```

[7, 8, 15]

Deep Mind plays Star Craft II

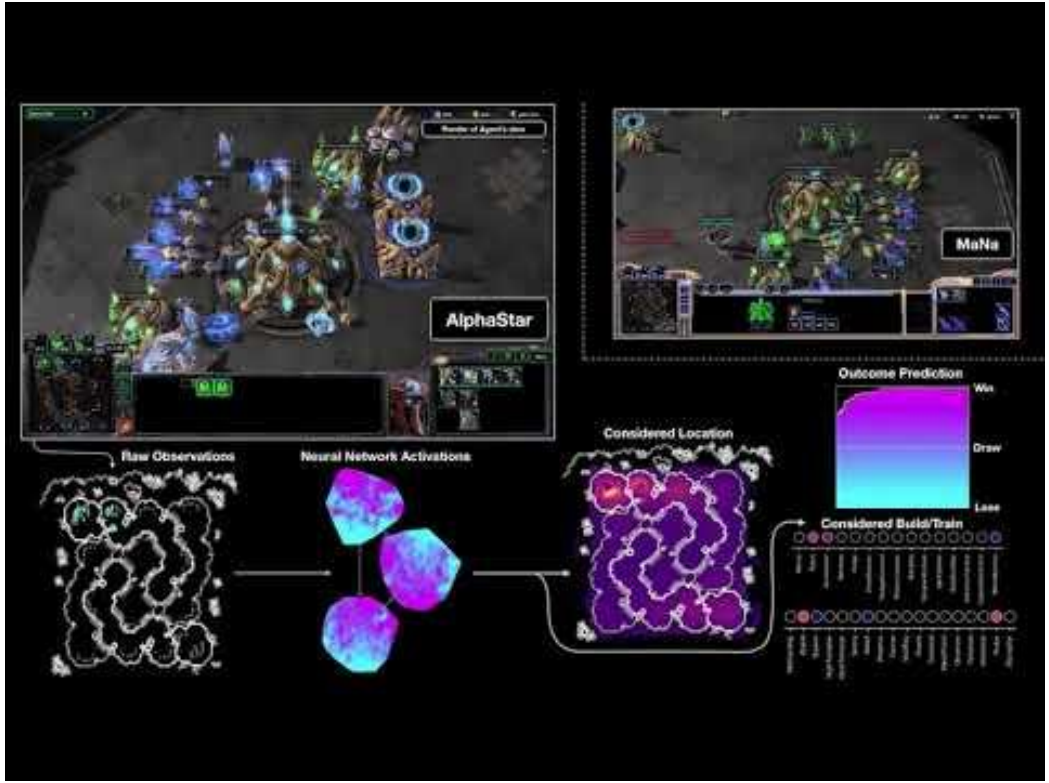

AlphaFold folds Proteins

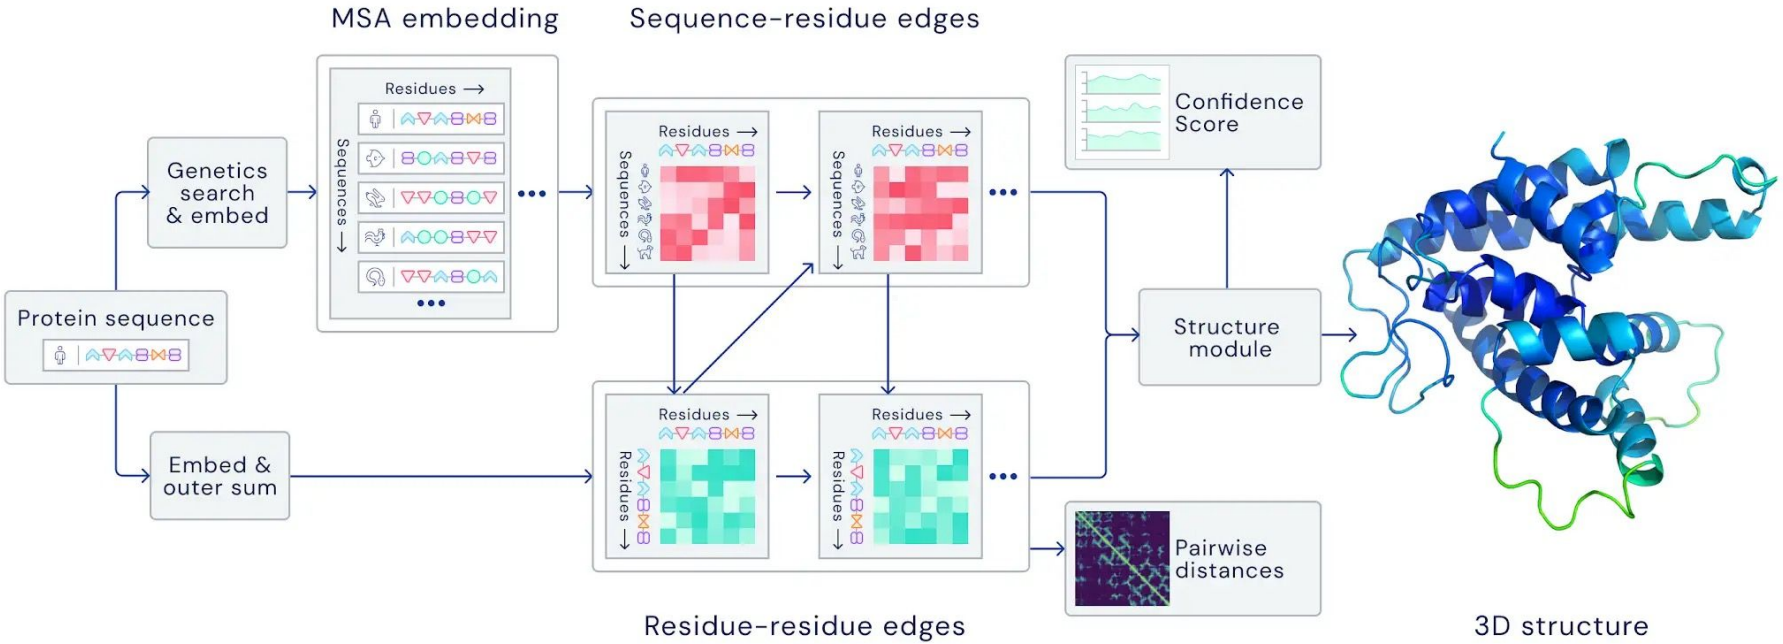

[17]

Summary

- Machine Learning can be an alternative to imperative programming
- Machine Learning = Using data to answer questions
- Deep Learning can do pretty cool stuff

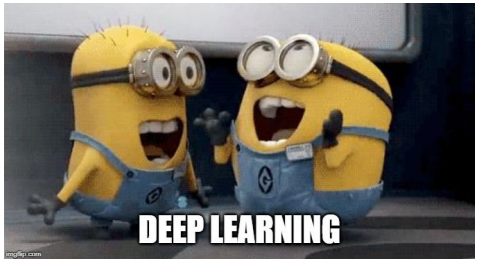

Literature

1. Ian Goodfellow and Yoshua Bengio and Aaron Courville. Deep Learning, 2016
2. Goyal et al. Artificial Intelligence for Diagnosis of Skin Cancer: Challenges and Opportunities., 2019
3. Vinyals et al. StarCraft II: A New Challenge for Reinforcement Learning, 2017
4. Ramesh et al. Zero-Shot Text-to-Image Generation, 2021
5. Lopez et al. Skin lesion classification from dermoscopic images using deep learning techniques, 2017
6. Goyal et al. Artificial Intelligence for Diagnosis of Skin Cancer: Challenges and Opportunities, 2019
7. Chen et al. Evaluating Large Language Models Trained on Code, 2021
8. Microsoft, OpenAI: OpenAI Model Generates Python Code, 2020
9. Leviathan et al. Google Duplex: An AI System for Accomplishing Real-World Tasks over the Phone, 2018
10. John Oliver. Robocalls: Last Week Tonight with John Oliver, 2019
11. Hadjeres et al. DeepBach: A Steerable Model for Bach Chorales Generation, 2017
12. Donahue et al. Piano Genie: An Intelligent Musical Interface, 2018
13. OpenAI. DALL-E: Creating images from text, 2021
14. OpenAI. DALL-E 2: Realistic images and art from a text description, 2022
15. Microsoft. GitHub Copilot: You AI Pair Programmer, 2023
16. DeepMind AlphaStar: Grandmaster level in StarCraft II and Demonstration, 2019
17. Jumper et al. Highly accurate protein structure prediction with AlphaFold, 2021
18. Veo3 video generator, 2025
19. Maslej et al. The AI Index 2025 Annual Report, 2025

Credits

Free icons from Flaticon:

- https://www.flaticon.com/free-icon/modeling_1373079
- https://www.flaticon.com/free-icon/analytics_166904
- https://www.flaticon.com/free-icon/landing_1284826

Other images:

Fig-1. Title image: https://flic.kr/p/28kDqNb
Fig-2. Cat image: https://bristolcountyvet.com/wp-content/uploads/2018/11/pexels-photo-min.jpg
Fig-3. Bagel vs. Muffin: https://www.boredpanda.com/dog-food-comparison-bagel-muffin-lookalike-teenybiscuit-karen-zack