

# Applied Deep Learning

---

Recurrent Neural Networks


Alexander Pacha - TU Wien

# Recap

- What do we mean when we say Multi Layer Perceptron (MLP)?
- What is the key insight that justifies Convolutional Layers?
- Why do we use parameter sharing and what are feature maps?




# Recap - Neural Network



# Modelling sequences with CNNs

CNNs expect grid-like structure and can operate on sequences

- Audio waveform, Color image data, Volumetric data, Color video data, ...



But:

- Fixed-size input
- Fixed-size output
- Fixed-size number of computations
- No understanding of previous information




We need to persist information

# Recurrent Neural Network Model


$A$  = a chunk of neural network

$x_t$  = Input at time step  $t$

$o_t$  = Output at time step  $t$




# Unfolding Computational Graph



# Parameter sharing

RNNs are specified in terms of *transition from one state to another*, with shared weights across all time steps, thus allowing to

- Generalize to sequence lengths not seen during training
- Share statistical strength across different positions in time
  - “I moved to Berlin in 2020” vs. “In 2020, I moved to Berlin”



# Backpropagation through time (BPTT)

- Simply apply back-propagation to unfolded graph
  - Can not parallelize because computation depends on previous state
  - Must step through entire graph
- Expensive to train: States must be stored until being reused in backward pass

Use of shared parameters in different time steps assumes that conditional probability distribution is stationary

→ Relationship between previous time step and next time step is independent of  $t$




# Activations


Most frequently used hyperbolic tangent function (tanh)

- Bounded output [-1; 1]
- Easier to compute than sigmoid


**Sigmoid**  
 $\sigma(x) = \frac{1}{1+e^{-x}}$




**tanh**  
 $\tanh(x)$



**ReLU**  
 $\max(0, x)$




**Leaky ReLU**  
 $\max(0.1x, x)$



**Maxout**  
 $\max(w_1^T x + b_1, w_2^T x + b_2)$

**ELU**  
$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Linear Units can have unbounded output and are likely to explode  
→ Avoid in RNN

# Vanishing/Exploding Gradients with RNNs

- Gradients propagated through many stages either vanish or explode
- Long-term dependencies must travel through many stages
  - Signal must be able to vanish
- Short-term dependencies have a much stronger signal
  - Small perturbations can interfere with signal from long-term dependencies


Therefore:

- The longer the span of dependencies that needs to be captured, the harder it is to train the network
- Special cells can mitigate this problem




# Gradient Clipping to avoid Exploding Gradients

Without clipping



With clipping



# Variations of RNNs


---



# RNN without outputs

Network just digests input into state  $\mathbf{h}$

■ = Delay of one time step



# RNN with single output at the end

$x$  = Input

$h$  = Hidden state

$y$  = Target

$o$  = Output

$L$  = Loss


$\tau$  = Number of steps

Weight matrices:

$U$  = input to hidden

$W$  = hidden to hidden

$V$  = hidden to output



# RNN with output at each time step

$x$  = Input

$h$  = Hidden state

$y$  = Target

$o$  = Output


$L$  = Loss

Weight matrices:


$U$  = input to hidden

$W$  = hidden to hidden

$V$  = hidden to output



# Teacher Forcing




# RNNs for Vector to Sequences

Unpacks input into sequence


Single vector input  $x$  used in

- initial state  $h_0$
- input at each time step
- both



# How to stop generation?

- Special symbol that can be generated to halt generation
  - Is added to each sequence from the training set
- Extra output (head) that predicts whether to stop or not
- Extra output that predicts sequence length  $\tau$ 
  - First predict  $\tau$  then sample  $\tau$  steps worth of data
  - Predict number of remaining steps  $\tau - t$



# RNN Sequence-to-Sequence of same length

$x$  = Input

$y$  = Target

$\circ$  = Output

$h$  = Hidden state

$L$  = Loss


Weight matrices:

$U$  = input to hidden

$W$  = hidden to hidden


$V$  = hidden to output

$R$  = output to hidden



# Bidirectional RNNs


- So far: Only took information from the past
- In some situations, the output depends on the whole sequence (e.g., speech recognition)
- Output units can benefit from past and future information






# Encoder-Decoder Sequence-to-Sequence

Input and output sequences can have different lengths:

- Process input with encoder (reader) into context C
- Decoder (writer) generates output from context C



# Deep RNNs




# Long Short-Term Memory (LSTM)

Designed to allow learning long-term dependencies


- Introduces self-loops to produce paths where gradient can flow for a long duration
- Introduces gates to change dynamically (controlled by another hidden unit)
- Use addition instead of multiplication

Gated Recurrent Unit (GRU) similar, but simpler (fewer gates)



# RNNs in Code

<https://colab.research.google.com/drive/1NXNcbzezSglkSUsG3pgv7yBwLAHG3dR0>




# Applications

---




# Short-Term Predictions

[https://colab.research.google.com/drive/1TIW\\_emPcv4YjjArDsgapZygOI\\_j3SF13](https://colab.research.google.com/drive/1TIW_emPcv4YjjArDsgapZygOI_j3SF13)




Source: <https://github.com/GoogleCloudPlatform/tensorflow-without-a-phd/tree/master/tensorflow-rnn-tutorial>

# CNNs and RNNs for Optical Character Recognition



Source: <https://github.com/apache/document-analysis/blob/master/assignment2>


# Recognizing Music Scores



Output: clef-G2, keySignature-DM, timeSignature-2/4, rest-sixteenth, note-F#4\_sixteenth, ...


Source [6]

# Image Captioning




Source: <https://ai.googleblog.com/2016/09/show-and-tell-image-captioning-open.html>

# Speech Recognition



Source: [7]


# Speech Recognition



# Connectionist Temporal Classification (CTC)



How to handle unaligned data?

*“My friend has too many cats”*



# Connectionist Temporal Classification (CTC)

Summing over probability of all possible alignments between input and the label.





Input	$x_1$	$x_2$	$x_3$	$x_4$	$x_5$	$x_6$	$x_7$	$x_8$	$x_9$	$x_{10}$	$x_{11}$	$x_{12}$
Alignment	h	h	a	s	—	t	t	o	ε	o	o	—
Output	h		a	s	—	t		o		o		—

Source: [9]



# Text translation




Spanish to English: [https://www.tensorflow.org/text/tutorials/nmt\\_with\\_attention](https://www.tensorflow.org/text/tutorials/nmt_with_attention)

# Modern Sequence Processing

- Due to the limitations of RNNs, sequence processing is nowadays mostly done with Attention and Transformers
- Fuels Large Language Models


→ Later lecture



# Summary

## Recurrent Neural Networks

- have feedback loops and can save state information
- are ideal to process sequential data where a memory is beneficial
- can be trained with Backpropagation Through Time (BPTT)
- are hard to train (vanishing / exploding gradient)
- can process arbitrary input and output sequences



## LSTMs / GRU

- can handle long-term dependencies
- replace multiplication with addition

CNNs and RNNs play really well together to solve challenging problems

# Literature

1. Deep Learning: <http://www.deeplearningbook.org/>
2. Karpathy on RNNs: <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>
3. Recurrent Neural Networks for Time Series Forecasting: Current Status and Future Directions:  
<https://arxiv.org/abs/1909.00590v1>
4. Simple OCR: <http://cs231n.github.io/assignments2019/assignment3/>
5. Understanding LSTMs (Blog): <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>
6. RNN-Walkthrough: <https://github.com/gabrielloye/RNN-walkthrough>
7. Neural End-To-End OMR: [http://ismir2018.ircam.fr/doc/pdfs/33\\_Paper.pdf](http://ismir2018.ircam.fr/doc/pdfs/33_Paper.pdf)
8. Recognizing Speech Commands:  
<https://towardsdatascience.com/recognizing-speech-commands-using-recurrent-neural-networks-with-attention-c2b2ba17c837>
9. Understanding CTC:  
<https://towardsdatascience.com/intuitively-understanding-connectionist-temporal-classification-3797e43a86c>
10. Neural Machine Translation with Attention mechanism: <https://arxiv.org/abs/1409.0473>
11. <https://towardsdatascience.com/a-comprehensive-guide-on-activation-functions-b45ed37a4fa5>
12. <https://machinelearningmastery.com/rectified-linear-activation-function-for-deep-learning-neural-networks/>

# Icon credits

Free icons from [Flaticon](#):

- [https://www.flaticon.com/free-icon/asking\\_900415](https://www.flaticon.com/free-icon/asking_900415)
- [https://www.flaticon.com/free-icon/balloons\\_609624](https://www.flaticon.com/free-icon/balloons_609624)
- [https://www.flaticon.com/free-icon/code\\_1383431](https://www.flaticon.com/free-icon/code_1383431)
- [https://www.flaticon.com/free-icon/crane\\_222566](https://www.flaticon.com/free-icon/crane_222566)
- [https://www.flaticon.com/free-icon/idea\\_427735](https://www.flaticon.com/free-icon/idea_427735)
- [https://www.flaticon.com/free-icon/transfer\\_179683](https://www.flaticon.com/free-icon/transfer_179683)
- [https://www.flaticon.com/free-icon/broken-link\\_1201519](https://www.flaticon.com/free-icon/broken-link_1201519)
- [https://www.flaticon.com/free-icon/gym\\_755336](https://www.flaticon.com/free-icon/gym_755336)
- [https://www.flaticon.com/free-icon/stop-sign\\_2168375](https://www.flaticon.com/free-icon/stop-sign_2168375)